


fasthade ch.1
by @dk21

DEEP LEARNING APPLICATIONS: NLP, VISION, MEDICINE, BIOLOGY, IMAGE GENERATION, RECOMMENDATION SYSTEMS, ROBOTICS, LOGISTICS, FINANCE



ML: discipline where we define program not by writing it entirely ourselves, but by learning from data

DL: subset of ML, using NN with multiple layers

TEACHING APPROACH: TOP-DOWN (VIA D. PERKINS, BASEBALL ANALOGY)

- 1) Start with real, practical examples (teach the whole game)
- 2) Learn by doing: run/re-implement code. Do personal projects
- 3) Deep dive into theory later, as needed - to improve models
- 4) Simplify and remove barriers: fastai library, Pytorch, Jupyter

LIMITATIONS OF ML:

- Need data with labels
- Can only learn patterns seen in ~~input~~ input data
- predictions vs recommended actions

WATCH OUT: FEEDBACK LOOPS!
e.g. Youtube recommending viral anti-vax videos

Some history:

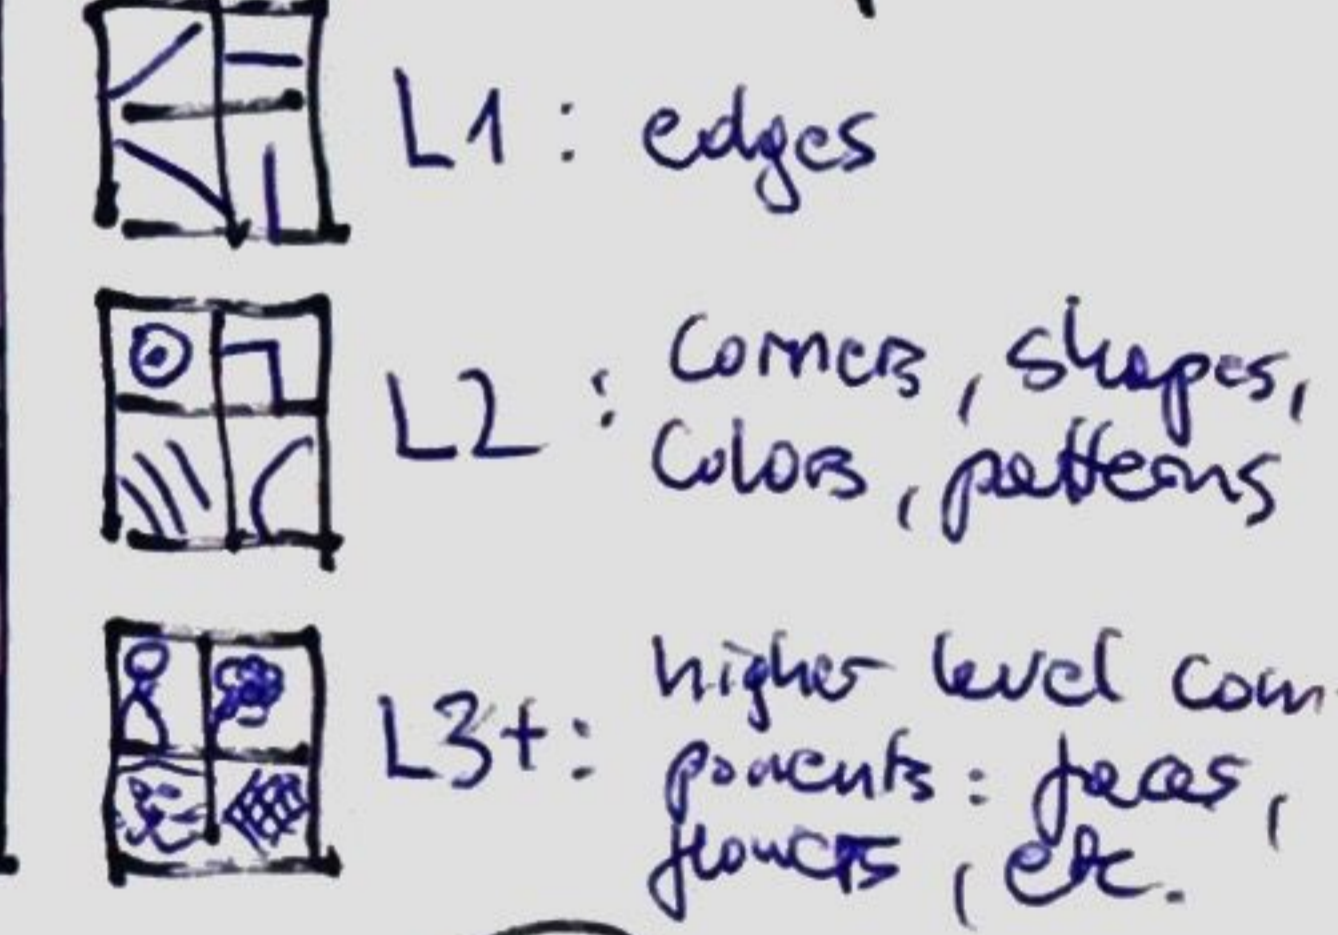
- 1943 - McCulloch, Pitts: Artificial Neuron
- Rosenblatt's device: Mark I Perceptron
- Minsky, Papert: Perceptrons - books introducing multiple layer neural networks
- 1986 - Parallel Distributed Processing - books introducing most of current DL framework
- 1961/62: Samuel - Artificial Intelligence essay, introducing current ML approach, program beating humans in checkers

```

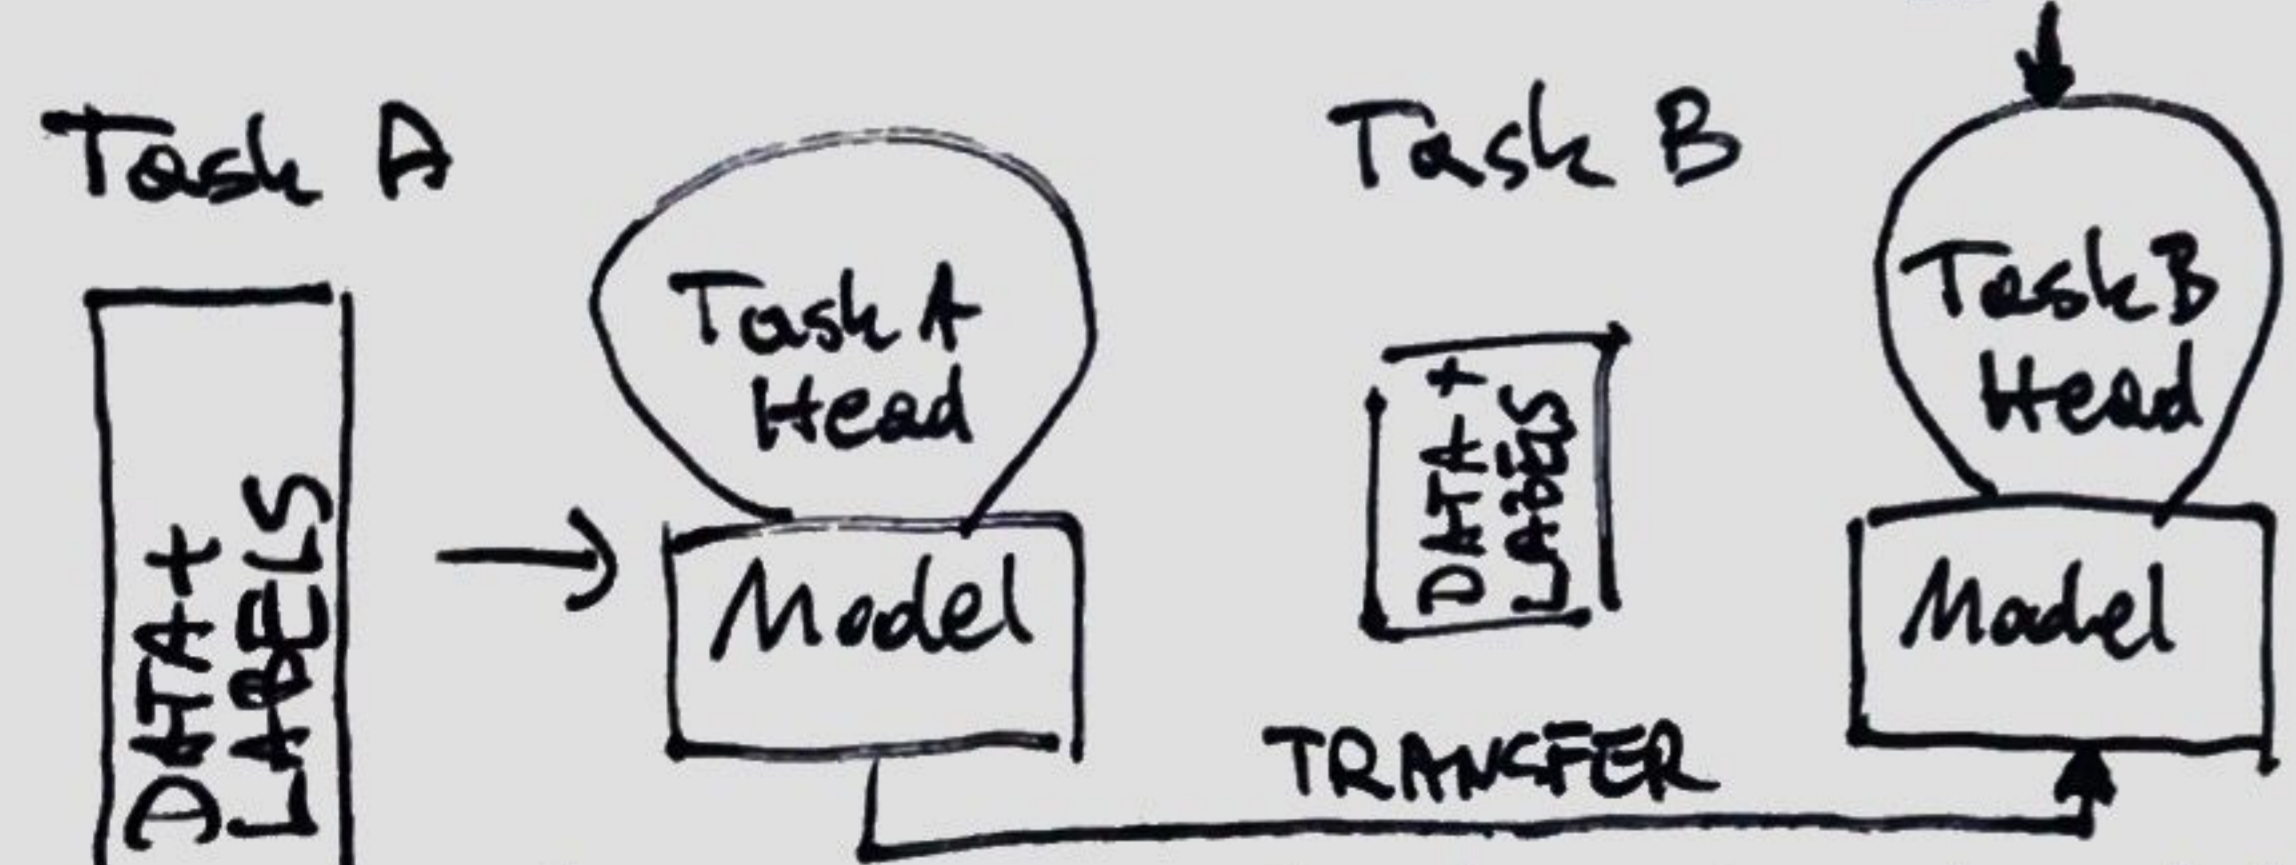
text | tabular | collab
from fastai.vision.all import * # fastai library
path = untar_data(URLs.PETS)/'images' # download dataset
def is_cat(x): return x[0].isupper() # labelling function
dls = ImageDataLoaders.from_name_func(
    path, get_image_files(path), valid_pct=0.2, seed=42,
    label_func=is_cat, item_tfms=Resize(224))
# load data, label, split into train/valid, transform
learn = cnn_learner(dls, resnet34, metrics=error_rate)
# load architecture, pretrained model, define metric
learn.fine_tune(1) # finetune ~ fit pretrained model
    
```

Different layers in NNs learn to recognize increasingly complex features.

CNN - Convolutional Neural Network example:




TRANSFER LEARNING:



- FINE-TUNE:
- 1) cut pre-trained model's head
 - 2) Add new head specific to new task
 - 3) Train new head only for 1 epoch
 - 4) Fit entire model for more epochs (more tricks here)

OVERFITTING: Single most important and challenging issue! Model starts memorizing examples, rather than learning to generalize!

As a rule, split data into train/validation set, or maybe test set as well:



AVOID LEAKAGE: Time series, same subjects in train/valid etc...

LOSS METRIC

Universal Approximation Theorem

CLASSIFICATION
REGRESSION
EPOCH PARAMETERS
ARCHITECTURE, CNN
SEGMENTATION
GPU, NOTEBOOK

OTHER TOPICS CONSTRAINED BY SPACE

fastbook ch 2 by @dlc21

DL STATE 2020 Computer Vision:
- object recognition
- object detection

Text: Classification, translation, summarization, Generation: compelling but not factually correct!
- question answering
- NER
- image captioning
- Other data types: eg protein chains

Tabular: timeseries, typically ensemble w/ RF, GBM helps: high-cardinality cols, eg ZIP.

Recommendation sys.: Collaborative filtering customers as rows products as columns

Underestimate DL capabilities

Overestimate DL capab.

DRIVETRAIN APPROACH: produce actionable outcomes, vs. predictions only

KEEP AN OPEN MIND

- 1 Complete lots of small experiments and work on your own project
2 Consider data availability
3 Iterate ELE - all the way to final product
4 Start with sth that DL is good at.

What outcome am I trying to achieve?

DEFINED OBJECTIVE

What inputs can we control?

LEVERS

What data can we collect?

DATA

How the levers influence the objective

MODELS

Randomized experiments needed to have good data, eg for recsys.

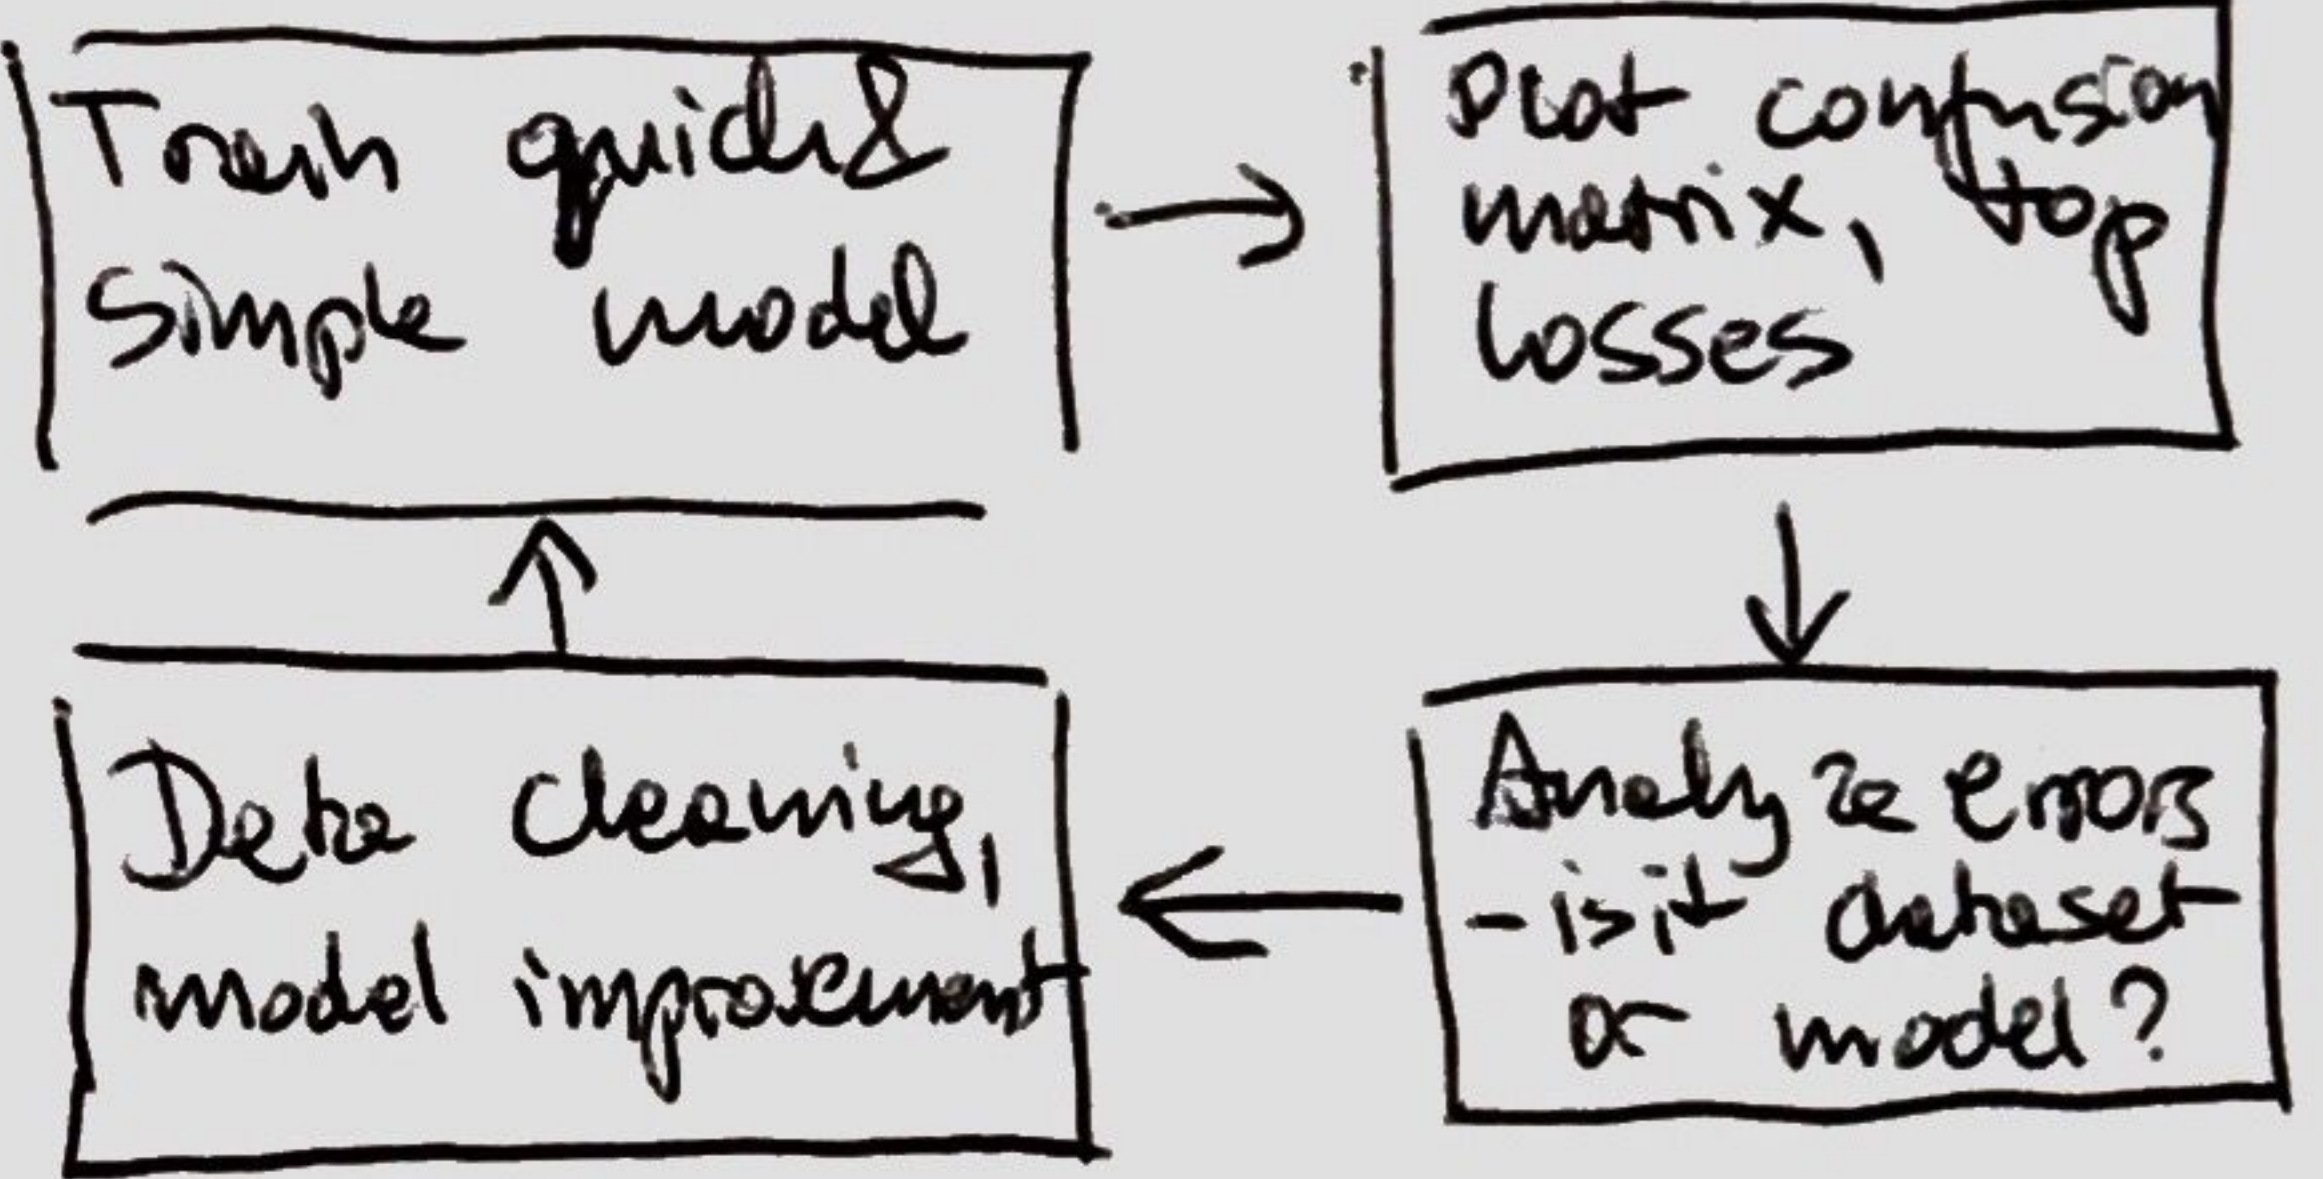
INFERENCE, CPU VS GPU

Voila ipywidgets

map method L. map (function)

(space constrained)

WORKFLOW:



```
class DataLoaders (getAttr): # provides data for model
def __init__(self, loaders): self.loaders = loaders
def __getitem__(self, i): return self.loaders[i]
train, valid = add_progs (lambda i, self: self[i])
```

```
bears = DataBlock ( # template for creating data loaders (DB)
blocks = (ImageBlock, CategoryBlock), # type of independent/dependent variable
get_items = get_image_files, # function takes a path and returns images in that path
splitter = RandomSplitter (valid_pct = 0.3, seed = 42), # split train/valid, fix seed
get_y = parent_label # label images
item_tfms = Resize (128) # item tfms run on CPU, eg. image resizing
dls = bears.dataloaders (path) # provide source of data to (DB) - here path with images
```

HOW TO AVOID DISASTER?

- Out of domain data: training vs production
• Domain shift: data changes over time
• Anticipate unforeseen consequences and feedback loops. What if this went really well?

RESIZE OPTIONS: crop, squish, pad, RandomResizedCrop & recommended

DATA AUGMENTATION: create random variation in the input data, standard set provided in aug-transforms, can be done on GPU in hkerl:

```
bears = bears.new (item_tfms = RandomResizedCrop (128, min_scale = 0.5), batch_tfms = aug_transforms)
Learn.export () : saves both model and parameters. load_learner (path/'export.pkl') = loads model (1)
```

GRADUAL ROLLOUT

1 Manual process

- Run model in parallel
- Humans check all predictions

2 Limited scope deployment

- Careful human supervision
- Time or geography limited

3 Gradual expansion

- Good reporting systems needed
- Consider what could go wrong!

TIP: Start writing, blog! write for people one step behind you.

ETHICS: the study of right & wrong,
DATA ETHICS: complicated, context

how we define & recognize them, understand the connection between action & consequences
dependent → learn through examples, like a muscle → develop & practice it!

BUGS, RECOURSE, ACCOUNTABILITY

- ⊗ Buggy algorithm cut healthcare benefits, impacting a vulnerable group
- ! Finger-pointing vs taking accountability
- ! bureaucracy as a way to evade responsibility
- ! Data often contains errors → mechanisms for audit and correction are crucial
- ⊗ Police maintaining database of gang members with no mechanism to correct obvious errors
- ⊗ US credit report system - very difficult to correct errors

FEEDBACK LOOPS

- ⊗ Conspiracy theories videos tend to get recommended more on YT | FB
- ! People watching them tend to watch more online videos
- ! YT | FB recommendation algorithms suggest more similar videos

WHY YOU SHOULD CARE?

- EX. IBM products used in Nazi concentration camps - would you be ok to contribute to killing people?
- EX. VW emission scandal - engineers failed!
- ! ML can create feedback loops & amplify bias
- ! People more likely to assume algorithms are objective and error-free
- ! Often used at scale, with no appeals process in place
- ! Considering this will make you a better practitioner!

BIAS Historical bias - people, processes, society are biased -

- ! taking real world data includes these biases
- MEDICAL - doctor prescriptions differ for white vs black patients
- SALES - different prices by race
- ⊗ Searching google for a name that is historically black, you get ads for background checks (suggesting a criminal record)
- ! Systematic imbalance in the make-up of popular datasets, eg. ImageNet, word embeddings
- ⊗ Translating doctor ~ man, nurse ~ woman.

Measurement bias - measuring wrong thing, in the wrong way, incorporating it into model incorrectly

- ⊗ Factors predictive of stroke
 - prior stroke
 - cardiovascular disease
 - accidental injury
 - colonoscopy
- ! these are correlated with people actually going to a doctor, being able to afford it, vs having a stroke

Aggregation bias - eg diabetes treatment based on linear, univariate models, small studies on homogeneous groups, when reality is non-linear, eg. diff. complications, symptoms across ethnicities

IDENTIFYING & ADDRESSING ETHICAL ISSUES

- 1 Analyze a project you're working on
 - Should we even be doing this?
 - What bias is in the data?
 - Can the code and data be audited?
 - What are error-rates for different sub-groups?
 - What is the accuracy of a simple, rule-based alternative?
 - What processes are in place to handle appeals or mistakes?
 - How diverse is the team that built it?
- 2 Processes to implement
 - EX. Regular ethical risk sweeps (pen testing)
 - include perspectives of a variety of stakeholders
 - what could bad actors do?
 - who will be directly and indirectly affected?
 - apply ethical lenses! which option...
 - [RIGHTS] best reflects the rights of all stakeholders
 - [JUSTICE] treat people equally or proportionately?
 - [UTILITARIAN] will produce most good & least harm?
 - [COMMON GOOD] best serves community as a whole, not just some members
 - [VIRTUE] leads me to act as the sort of person I want to be
- 3 The power of diversity
 - ! Similar backgrounds = similar blind spots
 - ! → Innovation, more risks/solutions considered
- 4 Role of policy - regulation is important
 - ⊗ FB lack of action during Rohingya genocide, vs quick action to address GDPR
 - ! Advocacy is important - support the regulations that you, a data scientist - believe we need!

spiral effects

`new_list = [f(o) for o in a_list if o > phi]`
list comprehension
 ↑
 do to do for each element, filter

TENSOR shape - length of each axis
 rank - number of axes = len(shape)

Measuring distance in space:
 Mean absolute distance (**L1 norm**)
 → absolute differences

Root mean squared error (**RMSE, L2 norm**)
 → mean of square diffs, then root

NUMPY ARRAY: multidimensional table of data, with all items of the same type, any type. With array of arrays, arrays underneath may have different sizes → "jagged array". Operations on regular arrays written in optimized C - much faster than Python.

PyTorch TENSOR - like a numpy array, but has to use simple basic numeric type for all components. Can run on GPU.

BROADCASTING - critical efficient code in PyTorch, when performing operation between tensors of different ranks, will automatically expand tensor with smaller rank to have the same size as the one with larger rank.

PyTorch VIEW change the shape of a tensor without changing contents, -1 parameter: make this axis as big as necessary to fit all data

@ - PyTorch matrix multiplication
 $batch @ weights + bias \Rightarrow$ fundamental operation MNN
 $W * X + b$
 ↑
 weights ↑
 biases } parameters
 Universal Approximation Theorem: this can represent any function

```
def train_model(model, epochs):
    for i in range(epochs):
        train_epoch(model)
        print(validate_epoch(model))
```

```
def train_epoch(model):
    for xb, yb in dl:
        calc_grad(xb, yb, model)
        opt.step()
        opt.zero_grad()
```

```
class BasicOptim:
    def __init__(self, params, lr):
        self.params, self.lr = list(params), lr
    def step(self, *args, **kwargs):
        for p in self.params:
            p.data -= p.grad.data * lr
# we use .data so PyTorch won't take gradient
# of this step
```

```
def zero_grad(self, *args, **kwargs):
    for p in self.params: p.grad = None
```

```
def calc_grad(xb, yb, model):
    preds = model(xb)
    loss = mnist_loss(preds, yb)
    loss.backward()
```

```
def linear(xb): return xb @ weights + bias
⇒ linear1 = nn.Linear(28*28, 1)
```

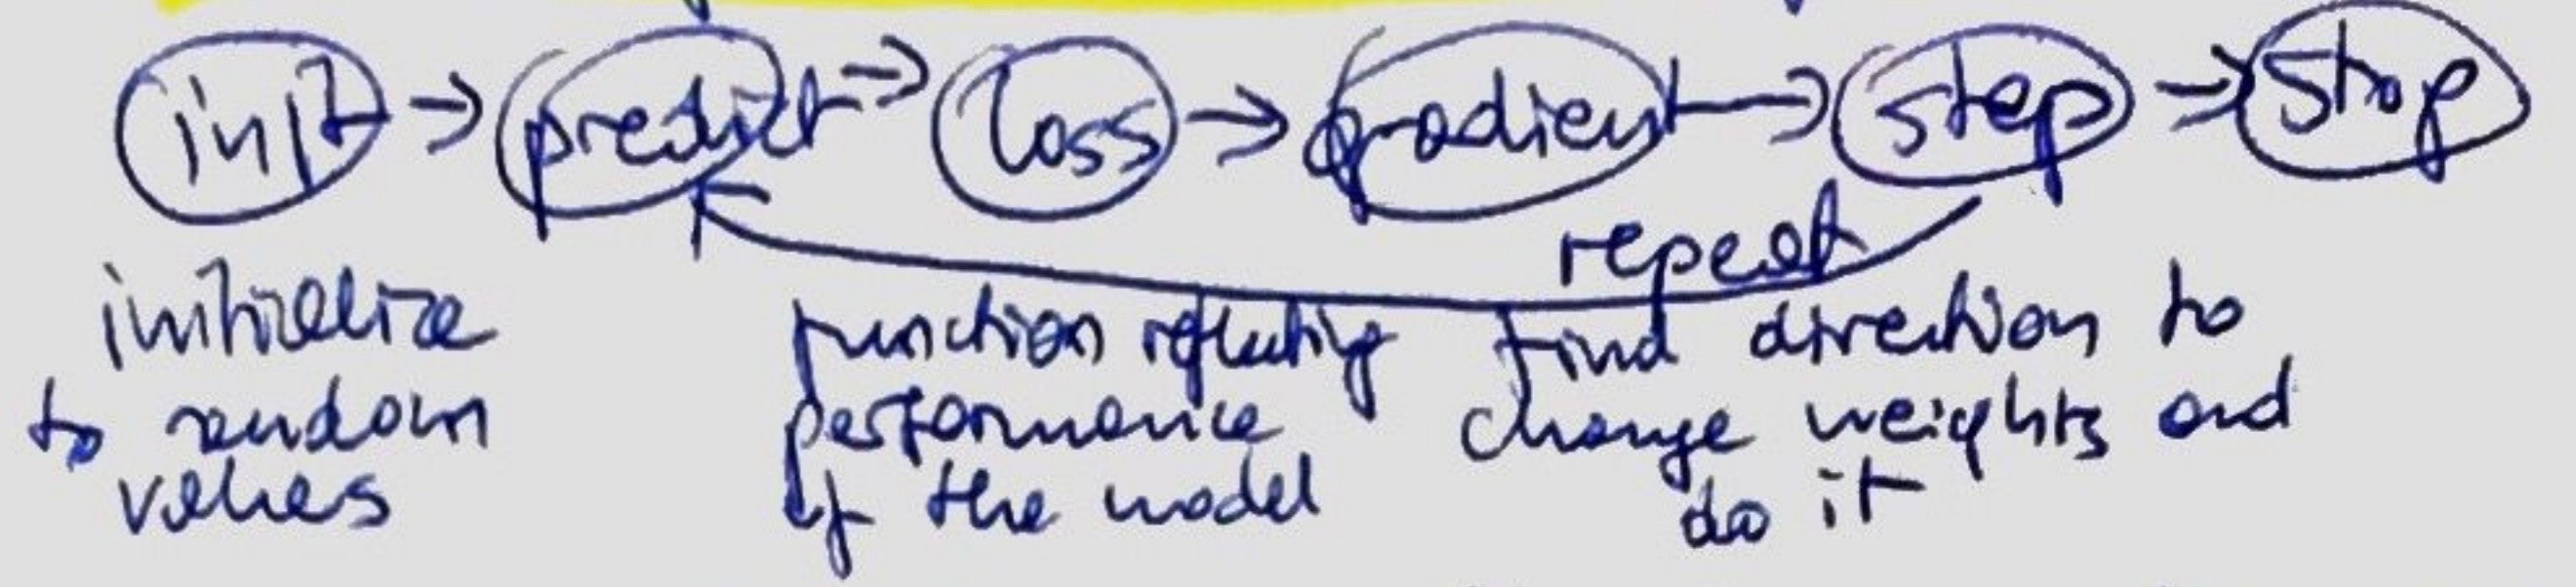
```
def mnist_loss(predictions, targets):
    predictions = predictions.sigmoid()
    return torch.where(targets == 1, 1 - predictions,
                       predictions).mean()
```

```
def init_params(size, std = 1):
    return torch.randn(size) * std, requires_grad_()
```

```
def simple_net(xb):
    res = xb @ w1 + b1
    res = res.max(tensor(0.0))
    res = res @ w2 + b2
    return res
```

← this is a neural network because we add a non-linearity (here: ReLU) between 2 classifiers
 in practice, we want more layers

SGD Gradient descent process:




Derivative of a function tells us how much a change in parameters will change results (rise/run)

GRADIENT: value of loss function's derivative at the point we're predicting. PyTorch can do it for us!
 def f(x): return x**2
 xt = tensor(3), requires_grad_()
 yt = f(xt)
 yt.backward()
 xt.grad # tensor(6.)



long time! explode vs converge!
 Multiply the gradient by a small learning rate to decide how much to change parameters.

LOSS FUNCTION: represents how good is the performance of our model. Needs to react to small changes in weights (accuracy isn't good!).



def sigmoid(x): return 1/(1+torch.exp(-x))
 → ensure values between 0 and 1.

METRIC - drive human understanding
Loss - drive automated learning.

To step ↓: change the weights/biases - we need to calculate loss on 1 or more data items. 1 is not enough - not much info, not optimized. Whole data set would be too slow → **MINI-BATCH**
 # items = batch size (! important decision)

We need to vary examples during training - randomly shuffle dataset on every epoch.

DATA LOADER: takes Python collection, turns it into iterator over batches:
 dl = DataLoader(collection, batch_size=8, shuffle=True)

DATA SET: collection with tuples of independent and dependent variables. Simplest PyTorch dataset.
 dataset = list(zip(x_train, y_train))

DEEP DIVE INTO MECHANICS OF DL

Learn Use eg in Regex! → RegexLabeler

RESIZING

```
item_tfms = Resize(460),
batch_tfms = aug_transforms(size=224, min_scale=0.75)
```

- 1) Resize images to relatively large dimension (vs target training dimension)
- 2) Compose all common aug operations (incl. resize to target size) into 1, and perform the combined operation on the GPU once at the end of processing
 - avoids data losses during augmentation
 - speeds up the process!

this adds RandomResizedCrop

CHECKING AND DEBUGGING DATA BLOCK

show_batch → inspect data, DataBlock.summary(path)

CROSS ENTROPY LOSS

Picture	Target	MODEL	LOGITS	SOFTMAX	NLL LOSS
#1	Teddy: 0	0	8 -5 -6	0.1 0.2 0.6	0.1 0.2 0.6
#2	Grizzly: 2	2	-4 6 3	0.05 0.95 0.05	0.05 0.95 0.05
#3	Brown: 1	1	-8 5 1	0.02 0.98 0.02	0.02 0.98 0.02

CROSS ENTROPY LOSS: 1.022

⇒ take the softmax, then negative log likelihood of that

Softmax: ensure final activations are between 0-1, and sum=1

```
def softmax(x): return exp(x)/exp(x).sum(dim=1, keepdim=True)
```


- if one activation is slightly higher than others, exp will amplify this - softmax really wants to pick one class - good if each image has definitive label
- may not be good at inference - will boost probs of example relative to other choices, independent of overall confidence - binary output columns, with sigmoid activation may be better?

Log Likelihood: pick loss from column with correct label only, take -log of that to transform 0↔1 scale to 0↔∞ scale

PyTorch: nn.CrossEntropyLoss ⇒ nn.LogSoftmax + nn.NLLLoss

Learning Rate Finder (Leslie Smith)

- start with a very small learning rate
- use that for 1 mini-batch, find the loss
- increase the lr gradually, eg double, per mini-batch, track the loss again
- keep doing this until the loss gets worse
- good choices: a) divide minimum loss lr by 10, or b) last point where the loss is clearly decreasing (steepest point)



Unfreezing and transfer learning

- remove pretrained model's classification head
- replace it with classif. head for new task
- this will have random weights initially, so we freeze pretrained layers and only train new head
- later, we unfreeze, check lr-finder again



Discriminative learning rates: pass slice (lr1, lr2) in faster

→ train first layers with smaller lr, last layer with higher lr, range between - multiplicatively equidistant lr's.

Selecting the number of epochs

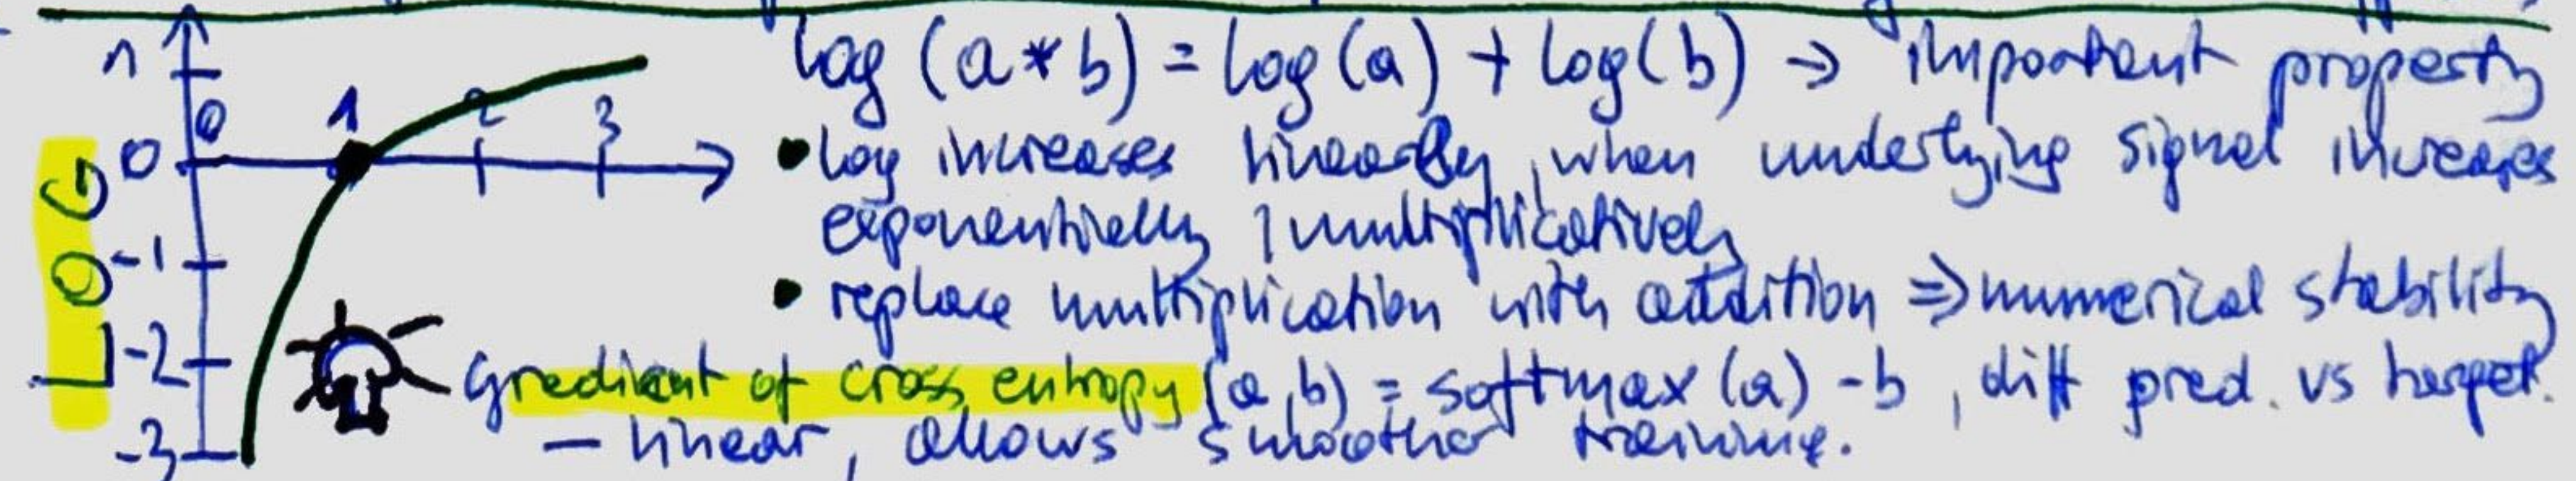
- 1) Choose based on time available
- 2) Observe training/validation losses and val. metrics
- 3) Make decisions on metrics, not losses! - initially, validation loss will get worse, because it becomes overconfident, it's still ok if the metrics improve. Only later, model will start to memorize.

Early stopping - save model after each epoch, select the one with best metric

This is NOT GOOD with 1 cycle training - epochs in the middle have higher lr, so unlikely to find best result. Better - if we overfit - retrain model from scratch with the number of epochs where we had best results.

Deeper Architectures: rule of thumb - more layers ⇒ more accurate model, but also risk of overfitting, longer to train, not always better!

We can speed this up with **mixed precision training**: learner.to_fp16()



Multi/one-hot encoding
[0 1 0 0 1 0 1 0 0]

BINARY CROSS ENTROPY (BCE)

def binary-cross-entropy(inputs, targets):
inputs = inputs.sigmoid()
return -torch.where(targets == 1, inputs, 1-inputs).log().mean()

Picture	Target	Logits	Sigmoid	Loss
#1	Teddy Brown Grizzly 1 0 0	8 -5 -6	1 0.01 0	0.001 0.
#2	0 1 0	-4 6 3	0.02 1 0.95	0.02 0.305
#3	0 1 1	-8 5 1	0 0.99 0.73	0 0.01 0.31

positive targets: $-\log(\text{sigmoid})$
neg. targets: $-\log(1-\text{sigmoid})$
BCE LOSS = 0.3777 (mean)

PyTorch: nn.BCELoss (without sigmoid) or nn.BCEWithLogitsLoss (with sigmoid)

loss_func = nn.BCEWithLogitsLoss()
loss = loss_func(activs, targets)

Accuracy - single label | Accuracy - multi label

```
def accuracy(inp, targ, axis=-1):
    pred = inp.argmax(dim=axis)
    return (pred == targ).float().mean()

def acc_multi(inp, targ, thresh=0.5, sigmoid=True):
    if sigmoid: inp = inp.sigmoid()
    return ((inp > thresh) == targ.bool()).float().mean()
```

PARTIAL example (Python):
acc_or = partial(acc_multi, thresh=0.2)

REGRESSION

Example - image regression, key point detection:
bim = DataBlock (blocks=(ImageBlock, PointsBlock),
get_items = get_image_files,
get_y = get_chr
splitter = FuncSplitter (lambda o: o.parent.name == 'B'),
batch_funcs = [* any-transforms (size=(240, 320)),
Normalize.from_stats (* imagenet_stats)]

Flexible API + Transfer Learning = POWER!

Rather than focusing on domains, focus on:

Independent var	Dependent var
Image	Text (caption)
Text	Image (from capt.)
Image + Text + Tabular	product - purchase-prob.

+ Loss function!

Finding the best threshold:
xs = torch.linspace(0.05, 0.95, 20)
accs = [acc_multi(pred, targ, thresh=i, sigmoid=False) for i in xs]

plt.plot(xs, accs);



ok to choose hyperparam based on valid set, if the function looks smooth

Multi-label classification: more than one type of object in an image - more common in practice to have some images with zero or more than 1 category match.

Use DataBlock API to construct DataLoader object from Pandas df:

→ start by creating and testing Datasets
→ create DataLoaders after that's working

dblock = DataBlock (get_x = lambda r: r['image'],
get_y = lambda r: r['labels'])

dsets = dblock.datasets(df)
dsets.train[0], len(dsets.train), len(dsets.valid)

! Lambda functions (defined inline) are good for iterating, but not compatible with serialization!
Need verbose functions to export Learner after training

def get_x(r): return path | 'train' | r['image']

def get_y(r): return r['labels'].split(' ')

def splitter(df):
train = df.index [~df ['is_valid']].tolist()
valid = df.index [df ['is_valid']].tolist()
return train, valid

dblock = DataBlock (blocks=(ImageBlock, MultiCategoryBlock),
splitter = splitter,
get_x = get_x,
get_y = get_y,
item_funcs = RandomResizedCrop (128, min_scale=0.35))

dl = dblock.dataloaders(df)

dl.show_batch (rows=1, ncols=3)

! Pass y-range to learner to force outputs into range: def sigmoid_range(x, lo, hi): return torch.sigmoid(x) * (hi-lo) + lo

learn = cnn_learner (dl, resnet18, y_range=(-1, 1))
dl.loss_func → MSELoss()
→ right function for regression!

- ① If your dataset is big, experiment on a subset of it
 - iterate at faster speed
 - the more experiments you can do, the better
 - subset should be representative → generalize

EX ImageNet: 10 classes from ImageNet

Normalization ⇒ mean = 0, stdev = 1

- helps the model train
- especially important when using pretrained models - distributed with stats used for normalization, use them for inference or transfer learning

```
check: x, y = dls.one_batch()
x.mean(dim=[0, 2, 3], x.std(dim=[0, 2, 3])
(average over all axes except channel=1)
```

fastai: add Normalize transform in batch_tfms

```
def get_dls(bs, size):
    dloader = DataBlock(blocks=(ImageBlock, CategoryBlock),
                        get_items=get_image_files,
                        get_y=parent_label,
                        item_tfms=Resize(460),
                        batch_tfms=[*aug_transforms(size=size,
                                                min_scale=0.75),
                                Normalize.from_stats(*imagenet_stats)])
    return dloader.dataloaders(path, bs=bs)
```

Progressive Resizing

Gradually using larger & larger images as we train similar to transfer learning: fine-tune after resizing works also as data augmentation

```
dls = get_dls(128, 128)
# create learner, fit_one_cycle
learn.dls = get_dls(64, 224)
# learn fine-tune (epochs, lr)
```

May hurt performance in transfer learning, if pretrained model/dataset are similar to our dataset.

Test Time Augmentation (TTA)

During inference or validation, creating multiple versions of each image using data augmentation, and then taking the average or maximum of the predictions.

- (the default in fastai is center-cropping for validation = largest square centered)
- problematic if relevant objects near edges
- squish/stretch may be difficult to train
- TTA solves these problems

```
preds, targs = learn.tta() # default is un-augmented center crop + 4 randomly augmented images, applied on valid dset.
```

Mixup

For each image img

- 1) select another dset image at random
- 2) pick a weight at random
- 3) take a weighted average of the img image and the selected image
- 4) take a weighted average of the img labels with the selected image's labels (targets need to be one-hot encoded)

In fastai: callbacks are used to inject custom behavior in training loop

```
model = xresnet50()
learn = Learner(dls, model, loss_func = CrossEntropyLossFlat(), metrics=accuracy, cbs = Mixup) # callback.
```

→ more epochs are needed to train for good accuracy (eg. ImageNet LB - mixup in models → 80 epochs)

⚠ can be used on activations inside models (NLP use cases)

check/read research papers

Label Smoothing (LS)

Problem with OHE: overconfidence, labels are always 0 or 1 even if there is nuance or uncertainty. Leads to overfitting, probabilities at inference not meaningful.

LS: replace all 1s with a number bit less than 1 - 0s - more than 0

- then train. Leads to:
 - training more robust, even with noisy data
 - models that generalize better

- ① Start with OHE - usually 0.1
- ② Replace all 0s with $\frac{\epsilon}{N}$ - no. of classes
- ③ Replace 1 with $1 - \epsilon + \frac{\epsilon}{N}$

In practice we don't change, or one-hot encode labels, but apply this in the loss function.

```
model = xresnet50()
learn = Learner(dls, model, loss_func = LabelSmoothingCrossEntropy(), metrics=accuracy, learn_fit_one_cycle(5, 3e-5))
→ more epochs are needed to train for good accuracy.
```

Summary

- ① Establish your environment for quick iteration (experimentation)
 - subset of dataset
 - validation generalizes to test/prod.
- ② Start with a simple, strong baseline
- ③ Run many experiments:
 - Augmentation
 - Lots of functions
 - Inspect data/results for insights

COLLABORATIVE FILTERING DEEP DIVE = look at what products the current user has used or liked, find other users that have used or liked similar products, then recommend other products that those users have used or liked. Generalize: items vs products, eg. diagnoses, links etc.

LATENT FACTORS

common characteristics of items / user preferences

Example:
 Movie is [0.9, 0.98, -0.9] A
 User likes [0.8, 0.5, -0.6] B

Dot Product A.B = Multiply vectors together, then sum up the result

We don't know latent factors → need to learn them!

Embedding from scratch in PyTorch

|| multiplying by a one-hot encoded matrix, using the computational shortcut that it can be implemented by simply indexing directly.

	idx	f1	f2	f3	
U1	1	2	3		← vector users3
U2	3	1	2		
U3	1	1	2		← matrix F
U4	2	2	3		

$F[3, :] \Rightarrow F^T * w$

We index into the embedding matrix using an integer, but calculate the derivative as if we were multiplying the matrix with OTE vector

Bootstrapping problem (cold start)

→ pick some user to represent avg taste
 → use a tabular model based on user metadata to construct initial embedding

! Representation bias & feedback loops are a risk - monitor keep humans in the loop, gradual rollout...

dls = Colab Data Loaders. from_dfl (ratings, item_name = 'title', bs=64)

n-users = len(dls.classes['user'])
 n-movies = len(dls.classes['title'])
 n-factors = 50

def create_params(size):
 return nn.Parameter(torch.zeros(*size).normal_(0, 0.01))

class DotProductBias (Module):
 def __init__(self, n-users, n-movies, n-factors, y-range=(0, 5.5)):

self.user-factors = create_params([n-users, n-factors])
 self.user-bias = create_params([n-users])
 self.movie-factors = create_params([n-movies, n-factors])
 self.movie-bias = create_params([n-movies])
 self.y-range = y-range

def forward(self, x):
 users = self.user-factors[x[:, 0]]
 movies = self.movie-factors[x[:, 1]]
 res = (users * movies).sum(dim=1)
 res += self.user-bias[x[:, 0]] + self.movie-bias[x[:, 1]]
 return sigmoid_range(res, *self.y-range)

PCA - principal component analysis

pull out underlying directions in latent factor matrix
 movie_w = learn.model.movie-factors[idxs].cpu().detach()
 movie_pca = movie_w.pca(3)
 fac0, fac1, fac2 = movie_pca.t()
 x = fac0[movie_idx]
 y = fac2[movie_idx]
 → pit scatter plot to visualize.

OOB in Python, **kwargs @delegates decorator NO SPACE

Weight decay / L2 regularization

Adding to our loss function the sum of all the weights squared, to encourage weights to stay small to prevent overfitting

wd | weight decay - parameter, how much we add loss:
 loss_with_wd = loss + wd * (parameters**2).sum()

parameters.grad += wd * 2 * parameters - same but faster!
 In torch, we pass it in a call to fit:
 learn.fit_one_cycle(5, 5e-3, wd=0.1)

Embedding distance

Movie similarity can be defined by the similarity of users that like those movies, distance between embedding vectors can define that similarity

movie-factors = learn.model.i_weight.weight
 idx = dls.classes['title'][0][2i['Movie title 1']]
 distances = nn.CosineSimilarity(dim=1)(movie-factors, movie-factors[idx][None])
 idx = distances.argsort(descending=True)[1]
 dls.classes['title'][idx]

Colab NN - dot product was PMF (probabilistic matrix factorization) approach, there is an option to do it with deep learning!

class ColabNN (Module):
 def __init__(...)
 ...
 self.layers = nn.Sequential(
 nn.Linear(user_sz[1] + item_sz[1], n_out),
 nn.ReLU(),
 nn.Linear(n_out, 1))
 ...
 def forward(self, x):
 embs = self.user-factors[x[:, 0]] self.item-factors[x[:, 1]]
 x = self.layers(torch.cat(embs, dim=1))
 return sigmoid_range(x, *self.y-range)

Fastback ch9 notes by @tk21
TABULAR MODELING ⇒ Data as a table; predict value in 1 column based on other cols

Variables

- **continuous**: numerical data feed directly to model
- **categorical**: discrete levels (eg movie ids), need to convert to numbers first
 - ↳ **Ordinal** - categories with natural ordering (eg 'ord-cat') Cat. set - categories (order, ordered=True, inplace=True)

CAT: represent via one-hot-encoding or **entity embedding**:
→ reduces memory usage and speeds up NN vs OHE
→ reveals intrinsic properties of variables - similar values close to each other in embedding space

Decision Trees TIP: avoid OHE categories for DTs | RFs

- 1) Loop through each column in dataset (greedy approach)
- 2) For each col, loop through each possible level of that col.
- 3) Try splitting data into 2 groups at that level
- 4) For regression: find avg value of dependent var. for each of 2 groups, see how close it is to the actual value of dep. variable for each of the items in that group
- 5) After looping thru all cols/levels, pick split point with the best predictions
- 6) For each group based on this split, repeat the process

Random Forests (Breiman 1984, 2001)

- 1) Randomly choose a subset of rows and subset of columns
- 2) Train a model using this subset (decision tree)
- 3) Save that model, return to step 1 several times
- 4) To make a prediction, predict using all saved models then take average of those predictions ⇒ **BAGGING**

↳ Important - errors of individual models are not correlated, so the average of those errors is ZERO

Out of Bag Error: measure prediction error on the training set by only including in the calculation of a row's error the rows where that row was not included in training

BOOSTING: another approach to ensemble (vs BAGGING)

- 1) Train a small model that underfits your dataset
- 2) Calculate predictions in the training set for this model
- 3) Subtract predictions from targets = **RESIDUALS**
- 4) Go back to step 1, now use the residuals as targets
- 5) Continue until a stopping criterion: max no trees, valid error getting worse, etc.

GBMs, GBDTs, XG Boost → Risk of overfitting
→ Very sensitive to hyperparameter choices

TABULAR MODELING WORKFLOW

- 1) Start with RF - easiest to train, resistant to hyperparam. choices, little preprocessing
- 2) Use RF model for feature selection, PDP analysis
- 3) Then try NNs or GBMs, try adding embeddings of cat. variables to the data

MODEL INTERPRETATION

- 1) How confident are we in our predictions for a particular row?
- 2) What were the most important factors influencing prediction?
- 3) Which columns are the strongest predictors, which can we ignore?
- 4) Which columns are effectively redundant?
- 5) How do predictions vary as we vary each column?

TREE VARIANCE

Check std deviation of predictions across trees:
preds = np.stack([t.predict(valid_xs) for t in m.estimators_])
preds_std = preds_std(0)
high std ⇒ low confidence

REMOVING LOW IMPORTANCE VARIABLES

Generally the 1st step to improve the model is simplifying it, so that it's easier to study, rollout, maintain

Remove columns/variables of low importance
→ Retrain the model → check impact on accuracy

REMOVING REDUNDANT FEATURES

cluster - columns (xs) - shows similarity of columns
We determine similarity by calculating the rank correlation - all values are replaced by their rank, then correlation is calculated
→ try removing each of potentially redundant features one at a time, then multiple from overlapping groups, observe CoB score / accuracy

TREE INTERPRETER + WATERFALL CHARTS

What were the most imp't factors for predicting with a particular row of data how did they influence that prediction? Calculated similarly to feature importances.
→ display feature contributions with waterfall chart
→ use it to provide useful information to users of your data product - reasons behind predictions

Extrapolation - Decision Trees / RF can never predict values outside of the range of training data (eg trend). NN can help. Also finding

Out-of-domain data: 1) combine train & valid datasets
2) Use RF to predict if a row is in train or valid set
3) get feature importances - for the columns that differ significantly between train/valid try removing them and see how it impacts accuracy. It can improve it, and make the model more resilient.
↳ Try to avoid using old data - it may no longer be predictive.

FEATURE IMPORTANCE

Loop through each tree, then recursively explore each branch - check what feature was used for that split, and how much the model improves as the result. The improvement, weighted by number of rows in that group, is added to importance score for that feature. Sum across all branches of all trees, normalize scores so that they add to 1.

PARTIAL DEPENDENCE PLOTS = PDP

PDP visualize how variables affect our predictions: if a row varied as nothing other than the feature in question, how would it impact the dependent variable?
- replace every value in Year Made col. with 1950
- calculate predicted sale price for every auction, take average over all auctions
- repeat for 1951, 1952, ..., 2020
- this isolates the effect of Year Made from sklearn, inspection
plot - partial dependence → The Book of why


DATA LEAKAGE

= giving model information about the target which normally should not be available at the time of prediction. How to detect it?
- check if the accuracy is too good to be true
- look for imp't predictors that don't make sense
- look for PDP plots that don't make sense in practice

Using a Neural Network

- 1) Decide which cols should be treated as cat vs cont
- 2) create embeddings for categorical variables
- 3) Add normalization (in proc)
- 4) Consider adding y-range to regression models
- 5) Adjust hidden layer sizes to size of dataset

fastai: TabularPandas (pandas df + convenience)
TabularProc tabular_learner other: dtreeviz
↳ Categorify TabularModel dtreeviz
↳ Fill Missing ↳ study source code add_datapart



ULMFit Approach: fine-tune pretrained LM on the target corpus prior to transfer learning to classification task.

Language Model: model trained to guess the next word in a text, after reading the words before

Self-supervised learning: training a model using labels that are embedded in the independent variable, rather than requiring external labels. Usually used for pretraining in transfer learning.

TEXT PREPROCESSING IN 3 STEPS

1) TOKENIZATION

= converting a text into a list of tokens

- a) **word-based**: split sentence on spaces and language-specific rules
- b) **subword-based**: analyze a corpus of documents to find the most commonly occurring groups of letters. These substrings become the vocab.

c) character-based

fastai adds some functionality with the `Tokenizer` class, e.g. special tokens:
`xxbos`: beginning of text
`xxmaj`: next word begins with a capital
`xxunk`: next word is unknown

see rules: `defaults.text_proc_rules`
`Setup` creates the vocab that is used in tokenization

```

spacy = WordTokenizer()
toks = first(spacy([txt]))
print(coll_repr(toks, 30))
tkn = Tokenizer(spacy)
print(coll_repr(tkn(txt), 31))
sp = SubwordTokenizer(vocab_sz = 52)
sp.setup(txts)
  
```

2) NUMERICALIZATION

= mapping tokens to integers

- 1) make a list of all possible levels of a variable (the vocab)
- 2) replace each level with its index in the vocab

```

num = Numericalize()
num.setup(toks)
coll_repr(num.vocabs, 20)
  
```

class methods

3) PUTTING TEXT INTO BATCHES

for a language model
 • we want LM to read text in order
 • we use a model that maintains a state - remembers what it read previously when predicting what comes next

- 1) transform indiv. texts into a stream by concatenating them, shuffle docs order before each epoch
- 2) cut this stream into a certain number of batches (batch size) mini-streams preserve order of tokens
- 3) Each time/step read `seq_len` from mini-streams

dependent variable is offset from the independent variable by 1 token // **LM Data Loader**

Potential to generate disinformation campaigns, flood social media with fake content etc -> see ch. 3 on ethics

collating items in a batch

- use padding to make texts all the same size
- sort (ish) docs by length prior to each epoch - batch together docs with similar lengths, pad to length of longest doc in a batch

```

get_imdb = partial(get_text_files, folders=['train', 'test', 'unsup'])
dls_lm = DataBlock(
  blocks = TextBlock.from_folder(path, is_lm=True),
  get_items = get_imdb, splitter = RandomSplitter(0.1)
).dataloaders(path, path=path, bs=128, seq_len=80)
dls_lm.show_batch(max_n=2)
  
```

when `TextBlock` passed `fastai` handles tokenization and numericalization

Words that are not in the vocab of pretrained LM will be added with random embeddings

fine-tuning language model

```

learn = language_model_learner(dls_lm, AND_LSTM,
  drop_mult=0.3, metrics=[accuracy, Perplexity()], to_fp16())
  
```

loss function: cross entropy
 perplexity metric: exponential of $-\text{torch.exp}(\text{cross_entropy})$

```

learn.fit_one_cycle(1, 2e-2)
learn.save('1epoch')
learn = learn.load('1epoch')
learn.unfreeze()
learn.fit_one_cycle(10, 2e-3)
learn.save_encoder('finetuned')
TEXT = 'I liked this movie because'
N_WORDS = 40
pred = learn.predict(TEXT, N_WORDS, temperature=0.75)
  
```

automatically frozen, this will only train embeddings
 use `fit_one_cycle` to save/load intermediate model results
ENCODER = model without task-specific final layers, like body in CNNs

TEXT GENERATION

creating the classifier dataloaders ! `is_lm=False`

```

dls_cls = DataBlock(
  blocks = (TextBlock.from_folder(path, vocab=dls_lm.vocab), CategoryBlock),
  get_y = parent_label,
  get_items = partial(get_text_files, folders=['train', 'test']),
  splitter = GrandparentSplitter(valid_name='test')
).dataloaders(path, path=path, bs=128, seq_len=72)
learn = text_classifier_learner(dls_cls, AND_LSTM, drop_mult=0.5,
  metrics=accuracy).to_fp16()
  
```

```

learn = learn.load_encoder('finetuned')
learn.fit_one_cycle(1, 2e-2)
learn.freeze_to(-2)
learn.fit_one_cycle(1, slice(1e-2/(2.6**4), 1e-2))
... (freeze_to(-3)) ...
learn.unfreeze()
learn.fit_one_cycle(2, slice(1e-3/(2.6**4), 1e-3))
  
```

gradual unfreezing + discriminative learning rates

Data Munging with PyTorch's Mid-Level API

Fastai is built on a layered API:
Top layer = applications - train a model in 5 lines of code, eg:
TextDataLoaders.from_folder()
Mid level API: - create new DataLoaders
- apply just part of transforms
- has the callback system, which allows to customize training loop any way we like
- has general optimizers

Transform: an object that behaves like a function, has optional setup method to initialize hidden state (eg vocabs) and an optional decode that reverses the function. Special behavior - always gets applied over tuples (input, target)

Examples: Tokenizer, Numericalize
1) Create object
2) call setup method
3) apply to input by calling object as a function
4) decode result back to understandable representation

Write your own Transform
1) Write a function / + decorator
def f(x: int): return x+1
f = Transform(f)
2) Decorator - Python syntax for passing a function to another function (or callable obj that behaves like a function)
@Transform
def f(x: int): return x+1
3) If we need setup or decode, then need to subclass Transform and implement encode

FROM: → → → →
path = untar_data(URLs.IMDB)
dls = DataBlock(blocks=(TextBlock.from_folder(path), CategoryBlock), get_y = parent_label, get_items = partial(get_text_files, folders=['train', 'test']), splitter = GrandParentSplitter(valid_name='test')).dataloaders(path)

class NormalizeMean(Transform)
def setup(self, items): self.mean = sum(items) / len(items)
def encode(self, x): return x - self.mean
def decode(self, x): return x + self.mean
tfm = NormalizeMean()
tfm.setup([1, 2, 3, 4, 5])

Pipeline: compose several Transforms together
tfms = Pipeline([tok, num]) define it by passing a list of transforms
t = tfms(txt) automatically applies transforms
tfms.decode(t) decode result of encoding

TfmdLists and Datasets: Transformed Collections
cut = int(len(files) * 0.8)
splits = [list(range(cut)), list(range(cut, len(files)))]
tlds = TfmdLists(files, [Tokenizer.from_folder(path), Numericalize], splits=splits)
t = tlds[0]
tlds.decode(t) data as a set of raw items (frames, dt rows...)
tlds.show(t) and Pipeline of Transforms. At initiation, will setup in each Transform in order. We make into it to get results of pipeline on row elements. Can handle train/valid: tlds.valid[0].

TO: (equivalent, but can be customized):
tfms = [[Tokenizer.from_folder(path), Numericalize], [parent_label, Categorize]]
files = get_text_files(path, folders=['train', 'test'])
splits = GrandParentSplitter(valid_name='test')(files)
dsets = Datasets(files, tfms, splits=splits)
dls = dsets.dataloaders(dl_type=SortedDL, before_batch=pad_input)

Datasets: apply 2 (or more) Pipelines in parallel to the same raw object and build a tuple with the result. → automatically do setup for us → when indexed, return a tuple with result of each pipeline.

x, y = dsets.valid[0]
x, y = dsets.valid[0]
tfms = [Tokenizer.from_folder(path), Numericalize]
y_tfms = [parent_label, Categorize()]
dsets = Datasets(files, [x_tfms, y_tfms], splits=splits)
x, y = dsets.valid[0]
→ convert it to dataloaders (here: need to pad input):
dls = dsets.dataloaders(bs=64, before_batch=pad_input)

DataLoader: collates the items from our datasets into batches. Impt ways to customize:
1) after_item: applied on each item after grabbing it inside the dataset (~ item_tfms in DataBlock)
2) before_batch: applied on list of items before they are collated. Ideal place to pad items to the same size.
3) after_batch: applied on the batch as a whole after its construction (~ batch_tfms)

Application Example: Siamese Pair
Siamese model takes two images and has to determine if they are same class or not
Application Example: Siamese Pair
Siamese model takes two images and has to determine if they are same class or not
Convert it to dataloaders w/ dataloader method